



ecruz0369 COLA and ARC files Python LAB4 

fafd123 · 2 minutes ago  History

# Debiasing word embeddings

Word embeddings are word vectors that have meaning, word vectors similar to each other will be close to each other in a vector space.

**After completing this lab you will be to:**

- Use and load pre-trained word vectors
- Measure similarity of word vectors using cosine similarity
- Solve word analogy problems such as Man is to Woman as Boy is to \_\_\_\_ using word embeddings
- Reduce gender bias in word embeddings by modifying word embeddings to remove gender stereotypes, such as the association between the words *receptionist* and *female*

## Word embeddings

Word embedding is a method used to represent words as vectors. They are popularly used in machine learning and natural language processing tasks. Despite their success in downstream tasks such as cyberbullying, sentiment analysis, and question retrieval, they exhibit gender stereotypes which raises concerns because their widespread use can amplify these biases.

Word embeddings are trained on word co-occurrence using a text dataset. After training, each word  $w$  will be represented as a  $d$ -dimensional word vector  $\vec{w} \in \mathbb{R}^d$ .

**Word embedding properties:**

- Words with similar semantic meaning will be close to each other
- The difference between word embedding vectors can represent relationships between words. For example, given the analogy "man is to King as woman is to  $x$ " (denoted as  $man : king :: woman : x$ ), by doing simple arithmetic on the embedding vectors, we find that  $x = queen$  is the best answer because  $\vec{man} - \vec{woman} \approx \vec{king} - \vec{queen}$ . For the analogy  $Paris : France :: Nairobi : x$ , finds that  $x = Kenya$ . These embeddings can also amplify sexism implicit in text. For instance,  $\vec{man} - \vec{woman} \approx \vec{computer\ programmer} - \vec{homemaker}$ . The same system that produced reasonable answers to the previous examples offensively answers "man is to computer programmer as woman is to  $x$ " with  $x = homemaker$ .

Run the following cell to load the required modules.

In [101...

```
import os
import json
import numpy as np
from pathlib import Path
from sklearn.decomposition import PCA
```

## Download and Load word vectors

Due to the computational resources required to train word embeddings, we will be using a pre-trained 50-dimensional word embeddings, GloVe to represent words.

Run the following cells to download and load the word embeddings.

In [102...

```
def download_glove_vectors():
    """
    Download the GloVe vectors
    Arguments:
        None
    Returns:
        file_name (String): The absolute path of the downloaded 50-dimensional
        GloVe word vector representations
    """

    if not Path('data').is_dir():
        print("Downloading the embeddings ...")
        !wget --quiet https://nlp.stanford.edu/data/glove.6B.zip
        print("Embeddings downloaded.")

        # Unzip it
        print("Unzipping the downloaded file ...")
        !unzip -q glove.6B.zip -d data/
        print("File unzipped.")

    return '/content/data/glove.6B.50d.txt'
```

In [103...

```
def get_glove_vectors(glove_file):
    """
    Read the word vectors in glove_file
    Arguments:
        glove_file (String): The absolute path to the downloaded glove word embeddings
    Returns:
        words (Set): The words (vocabulary) in the pretrained glove word embeddings
        word_to_vector_map (Dict): A dictionary mapping the each word to its embedding vector
    """

    words = set()
    word_to_vector_map = {}
    with open(glove_file, 'r') as file_handle:
        for line in file_handle:
            line = line.strip().split()
            current_word = line[0]
```

```

        current_word = line[0]
        words.add(current_word)
        current_word_vector = line[1:]
        word_to_vector_map[current_word] = np.array(current_word_vector, dtype=np.float64)

    return words, word_to_vector_map

```

In [104...

```

# Load sets of words in the vocabulary and a dictionary mapping words to their GloVe vectors
words, word_to_vector_map = get_glove_vectors(download_glove_vectors())

```

## Operations on word embeddings

### Task 1 - Cosine similarity

Similarity between two words represented as word vectors  $u$  and  $v$  can be measured by their cosine similarity:

$$\text{CosineSimilarity}(u, v) = \frac{u \cdot v}{\|u\|_2 \|v\|_2} = \cos(\theta) \quad (1)$$

Where:

$u \cdot v$  is the dot (inner) product of the two vectors

$\|u\|_2$  is the length of the vector  $u$ . The length also called Euclidean length or Euclidean norm defines a distance function defined as

$$\|u\|_2 = \sqrt{u_1^2 + \dots + u_n^2}$$

The normalized similarity between  $u$  and  $v$  is the cosine of the angle between the two vectors denoted as  $\theta$ . The cosine similarity of  $u$  and  $v$  will be close to 1 if the two vectors are similar, otherwise, the cosine similarity will be small.

**Note:** We will be referring to the embedding of a word i.e the word vector and the word interchangeably in this lab.

---

**Task 1a:** Implement equation 1 in the `cosine_similarity()` function below.

Hint: check out the numpy documentation on [np.dot](#), [np.sum](#), and [np.sqrt](#). Depending on how you choose to implement it, you can check out [np.linalg.norm](#).

---

In [105...

```

def cosine_similarity(vector1, vector2):
    """
    Calculates the cosine similarity of two word vectors - vector1 and vector2
    Arguments:
        vector1 (ndarray): A word vector having shape (n,)
        vector2 (ndarray): A word vector having shape (n,)
    Returns:

```

```

Returns:
    cosine_similarity (float): The cosine similarity between vector1 and vector2
    """

    # Compute the dot product between vector1 and vector2
    dot = np.dot(vector1, vector2)

    # Compute the Euclidean norm or Length of vector1
    norm_vector1 = np.linalg.norm(vector1)

    # Compute the Euclidean norm or Length of vector2
    norm_vector2 = np.linalg.norm(vector2)

    # Compute the cosine similarity as defined in equation 1
    cosine_similarity = dot / (norm_vector1 * norm_vector2)

    return cosine_similarity

```

In [106...

```

# Run this cell to obtain and report your answers
man = word_to_vector_map["man"]
woman = word_to_vector_map["woman"]
cat = word_to_vector_map["cat"]
dog = word_to_vector_map["dog"]
orange = word_to_vector_map["orange"]
england = word_to_vector_map["england"]
london = word_to_vector_map["london"]
edinburgh = word_to_vector_map["edinburgh"]
scotland = word_to_vector_map["scotland"]

print(f"Cosine similarity between man and woman: {cosine_similarity(man, woman)}")
print(f"Cosine similarity between cat and dog: {cosine_similarity(cat, dog)}")
print(f"Cosine similarity between cat and cow: {cosine_similarity(cat, orange)}")
print(f"Cosine similarity between england - london and edinburgh - scotland: {cosine_similarity(england - london, edinburgh -

```

```

Cosine similarity between man and woman: 0.886033771849582
Cosine similarity between cat and dog: 0.9218005273769252
Cosine similarity between cat and cow: 0.40695688711826294
Cosine similarity between england - london and edinburgh - scotland: -0.5203389719861108

```

**Task 1b:** In the code cell below, try out 3 of your own inputs here and report your inputs and outputs

In [107...

```

# Run this cell to obtain and report your answers
king = word_to_vector_map["king"]
queen = word_to_vector_map["queen"]
tiger = word_to_vector_map["tiger"]
lion = word_to_vector_map["lion"]
apple = word_to_vector_map["apple"]
france = word_to_vector_map["france"]

```

```

paris = word_to_vector_map["paris"]
berlin = word_to_vector_map["berlin"]
germany = word_to_vector_map["germany"]

print(f"Cosine similarity between king and queen: {cosine_similarity(king, queen)}")
print(f"Cosine similarity between tiger and lion: {cosine_similarity(tiger, lion)}")
print(f"Cosine similarity between tiger and apple: {cosine_similarity(tiger, apple)}")
print(f"Cosine similarity between france - paris and berlin - germany: {cosine_similarity(france - paris, berlin - germany)}")

```

```

Cosine similarity between king and queen: 0.7839043010964117
Cosine similarity between tiger and lion: 0.5493991225945609
Cosine similarity between tiger and apple: 0.36045373632604055
Cosine similarity between france - paris and berlin - germany: -0.795694872209922

```

## Task 2 - Word analogy

In an analogy task, you are given an analogy in the form "i is to j as k is to \_\_\_\_". Your task is to complete this sentence.

For example, if you are given "man is to king as woman is to *l*" (denoted as  $man : king :: woman : l$ ). You are to find the best word *l* that answers the analogy the best. Simple arithmetic of the embedding vectors will find that  $l = queen$  is the best answer because the embedding vectors of words *i*, *j*, *k*, and *l* denoted as  $e_i$ ,  $e_j$ ,  $e_k$ ,  $e_l$  have the following relationship:

$$e_j - e_i \approx e_l - e_k$$

Cosine similarity can be used to measure the similarity between  $e_j - e_i$  and  $e_l - e_k$

**Task 2a:** To perform word analogies, implement `answer_analogy()` below.

In [108...

```

# Fixed the issue this code now works Nov 14, 2024 ECruz

def answer_analogy(word_i, word_j, word_k, word_to_vector_map):
    """
    Performs word analogy as described above
    Arguments:
        word_i (String): A word
        word_j (String): A word
        word_k (String): A word
        word_to_vector_map (Dict): A dictionary of words as key and its associated embedding vector as value
    Returns:
        best_word (String): A word that fulfils the relationship that e_j - e_i as close as possible to e_l - e_k, as measured
    """

    # Convert words to lowercase
    word_i = word_i.lower()

```

```

word_j = word_j.lower()
word_k = word_k.lower()

# Start code here #
try:
    # Get the embedding vectors of word_i (~ 1 Line)
    embedding_vector_of_word_i = word_to_vector_map[word_i] # Fetch embedding from word_to_vector_map
except KeyError:
    print(f"{word_i} is not in our vocabulary. Please try a different word.")
    return

try:
    # Get the embedding vectors of word_j (~ 1 Line)
    embedding_vector_of_word_j = word_to_vector_map[word_j] # Fetch embedding from word_to_vector_map
except KeyError:
    print(f"{word_j} is not in our vocabulary. Please try a different word.")
    return

try:
    # Get the embedding vectors of word_k (~ 1 Line)
    embedding_vector_of_word_k = word_to_vector_map[word_k] # Fetch embedding from word_to_vector_map
except KeyError:
    print(f"{word_k} is not in our vocabulary. Please try a different word.")
    return

# End code here #

# Get all the words in our word to vector map i.e our vocabulary
words = word_to_vector_map.keys()
max_cosine_similarity = -1000 # Initialize to a large negative number
best_word = None # Note: Do not change this None. Keeps track of the word that best

# Since we are looping through the whole vocabulary, if we encounter a word
# that is the same as our input, that word becomes the best_word. To avoid
# that we skip the input word.
input_words = set([word_i, word_j, word_k])

for word in words:
    if word in input_words:
        continue

    # Start code here #
    # Compute cosine similarity (~ 1 Line)
    # Calculate cosine similarity using embedding vectors
    embedding_vector_of_word_l = word_to_vector_map[word]
    a = embedding_vector_of_word_j - embedding_vector_of_word_i
    b = embedding_vector_of_word_l - embedding_vector_of_word_k
    similarity = np.dot

```

**Task 2b:** Test your implementation by running the code cell below. What are your observations? What do you observe about the last two outputs?.

In [109...

```
analogies = [('france', 'french', 'germany'),
              ('england', 'london', 'japan'),
              ('boy', 'girl', 'man'),
              ('man', 'doctor', 'woman'),
              ('small', 'smaller', 'big')]
for analogy in analogies:
    best_word = answer_analogy(*analogy, word_to_vector_map)
    if best_word:
        print(f"{analogy[0]} -> {analogy[1]} :: {analogy[2]} -> {best_word}")
```

**Task 2c:** Try your own analogies by completing and executing the code cell below. Find 2 that works and one that doesn't. Report your inputs and outputs

In [110...

```
# Define your own analogies here
my_analogies = [
    ('king', 'queen', 'prince'),
    ('paris', 'france', 'nice'),
    ('car', 'road', 'sign')
]

# Execute the analogies
for analogy in my_analogies:
    best_word = answer_analogy(*analogy, word_to_vector_map)
    print(f"{analogy[0]} -> {analogy[1]} :: {analogy[2]} -> {best_word}")
```

```
king -> queen :: prince -> None
paris -> france :: nice -> None
car -> road :: sign -> None
```

## Task 3 - Geometry of Gender and Bias in Word Embeddings: Occupational stereotypes

In this task, we will understand the biases present in word-embedding i.e which words are closer to *she* than to *he*. This will be achieved by evaluating whether the GloVe embeddings have stereotypes on occupation words. Determine gender bias by projecting each of the occupations onto the *she* — *he* direction by computing the dot product between each occupation word embedding and the embedding vector of *she* — *he* normalized by the Euclidean norm (See task 1).



$$occupation\_word_i \cdot ||she - he||_2$$

(2)

Notice that equation 2 is similar to only the numerator of equation 1 because we are computing the dot product of  $occupation\_word_i$  and the normalized difference between  $she$  and  $he$ .

Run the cells below to download and view the occupations.

In [111...

```
def download_occupations():
    if not Path('debiaswe').is_dir():
        print("Downloading occupation list ...")
        !git clone -q https://github.com/tolga-b/debiaswe.git
        print("Occupation list downloaded.")

    return '/content/debiaswe/data/professions.json'

def view_occupations(occupations_file):
    with open(occupations_file, 'r') as file_handle:
        occupations = json.load(file_handle)

        for occupation in occupations:
            print(occupation[0])
```

In [112...

```
occupations_file = download_occupations()
```

In [113...

```
view_occupations(occupations_file)
```

```
accountant
acquaintance
actor
actress
adjunct_professor
administrator
adventurer
advocate
aide
alderman
alter_ego
ambassador
analyst
anthropologist
archaeologist
archbishop
architect
artist
artiste
assassin
assistant_professor
```

assistant\_professor  
associate\_dean  
associate\_professor  
astronaut  
astronomer  
athlete  
athletic\_director  
attorney  
author  
baker  
ballerina  
ballplayer  
banker  
barber  
baron  
barrister  
bartender  
biologist  
bishop  
bodyguard  
bookkeeper  
boss  
boxer  
broadcaster  
broker  
bureaucrat  
businessman  
businesswoman  
butcher  
butler  
cab\_driver  
cabbie  
cameraman  
campaigner  
captain  
cardiologist  
caretaker  
carpenter  
cartoonist  
cellist  
chancellor  
chaplain  
character  
chef  
chemist  
choreographer  
cinematographer  
citizen  
civil\_servant  
cleric  
clerk  
coach  
11

collector  
colonel  
columnist  
comedian  
comic  
commander  
commentator  
commissioner  
composer  
conductor  
confesses  
congressman  
constable  
consultant  
cop  
correspondent  
councilman  
councilor  
counselor  
critic  
crooner  
crusader  
curator  
custodian  
dad  
dancer  
dean  
dentist  
deputy  
dermatologist  
detective  
diplomat  
director  
disc\_jockey  
doctor  
doctoral\_student  
drug\_addict  
drummer  
economics\_professor  
economist  
editor  
educator  
electrician  
employee  
entertainer  
entrepreneur  
environmentalist  
envoy  
epidemiologist  
evangelist  
farmer  
fashion\_designer

fighter\_pilot  
filmmaker  
financier  
firebrand  
firefighter  
fireman  
fisherman  
footballer  
foreman  
freelance\_writer  
gangster  
gardener  
geologist  
goalkeeper  
graphic\_designer  
guidance\_counselor  
guitarist  
hairstylist  
handyman  
headmaster  
historian  
hitman  
homemaker  
hooker  
housekeeper  
housewife  
illustrator  
industrialist  
infielder  
inspector  
instructor  
interior\_designer  
inventor  
investigator  
investment\_banker  
janitor  
jeweler  
journalist  
judge  
jurist  
laborer  
landlord  
lawmaker  
lawyer  
lecturer  
legislator  
librarian  
lieutenant  
lifeguard  
lyricist  
maestro  
magician

magistrate  
maid  
major\_leaguer  
manager  
marksman  
marshal  
mathematician  
mechanic  
mediator  
medic  
midfielder  
minister  
missionary  
mobster  
monk  
musician  
nanny  
narrator  
naturalist  
negotiator  
neurologist  
neurosurgeon  
novelist  
nun  
nurse  
observer  
officer  
organist  
painter  
paralegal  
parishioner  
parliamentarian  
pastor  
pathologist  
patrolman  
pediatrician  
performer  
pharmacist  
philanthropist  
philosopher  
photographer  
photojournalist  
physician  
physicist  
pianist  
planner  
plastic\_surgeon  
playwright  
plumber  
poet  
policeman  
politician

pollster  
preacher  
president  
priest  
principal  
prisoner  
professor  
professor\_emeritus  
programmer  
promoter  
proprietor  
prosecutor  
protagonist  
protege  
protester  
provost  
psychiatrist  
psychologist  
publicist  
pundit  
rabbi  
radiologist  
ranger  
realtor  
receptionist  
registered\_nurse  
researcher  
restaurateur  
sailor  
saint  
salesman  
saxophonist  
scholar  
scientist  
screenwriter  
sculptor  
secretary  
senator  
sergeant  
servant  
serviceman  
sheriff\_deputy  
shopkeeper  
singer  
singer\_songwriter  
skipper  
socialite  
sociologist  
soft\_spoken  
soldier  
solicitor  
solicitor\_general

solicitor\_general  
soloist  
sportsman  
sportswriter  
statesman  
steward  
stockbroker  
strategist  
student  
stylist  
substitute  
superintendent  
surgeon  
surveyor  
swimmer  
taxi\_driver  
teacher  
technician  
teenager  
therapist  
trader  
treasurer  
trooper  
trucker  
trumpeter  
tutor  
tycoon  
undersecretary  
understudy  
valedictorian  
vice\_chancellor  
violinist  
vocalist  
waiter  
waitress  
warden  
warrior  
welder  
worker  
wrestler  
writer

---

**Task 3a:** Complete the `get_occupation_stereotypes()` below.

---

In [114...

```
#completed code Edward Cruz
```

```
def get_occupation_stereotypes(she, he, occupations_file, word_to_vector_map, verbose=False):
```

```
    """
```

```
    Computes the words that are closest to she and he in the GloVe embeddings
```

computes the words that are closest to she and he in the GloVe embeddings

Arguments:

she (String): A word

he (String): A word

occupations\_file (String): The path to the occupation file

word\_to\_vector\_map (Dict): A dictionary mapping words to embedding vectors

Returns:

most\_similar\_words (Tuple(List[Tuple(Float, String)], List[Tuple(Float, String)])):

A tuple of the list of the most similar occupation words to she and he with their associated similarity

"""

*# Read occupations*

with open(occupations\_file, 'r') as file\_handle:

occupations = json.load(file\_handle)

*# Extract occupation words*

occupation\_words = [occupation[0] for occupation in occupations]

*# Get embedding vector of she*

embedding\_vector\_she = word\_to\_vector\_map[she]

*# Get embedding vector of he*

embedding\_vector\_he = word\_to\_vector\_map[he]

*# Get the vector difference between embedding vectors of she and he*

vector\_difference\_she\_he = embedding\_vector\_she - embedding\_vector\_he

*# Get the normalized difference*

normalized\_difference\_she\_he = vector\_difference\_she\_he / np.linalg.norm(vector\_difference\_she\_he)

*# Store the cosine similarities*

similarities = []

for word in occupation\_words:

try:

*# Get the embedding vector of the current occupation word*

occupation\_word\_embedding\_vector = word\_to\_vector\_map[word]

*# Compute cosine similarity between embedding vector of the occupation word and normalized she - he vector*

similarity = np.dot(occupation\_word\_embedding\_vector, normalized\_difference\_she\_he) / np.linalg.norm(occupation\_w

similarities.append((similarity, word))

except KeyError:

if verbose:

print(f"{word} is not in our vocabulary.")

*# Sort similarities*

most\_similar\_words = sorted(similarities)

return most\_similar\_words[:20], most\_similar\_words[-20:]



**Task 3b:** Execute the cell below and report your results.

1. Does the GloVe word embeddings propagate bias? why?
2. From the list associated with she, list those that reflect gender stereotype.
3. Compare your list from 2 to the occupations closest to he. What are your conclusions?

Exclude businesswoman from your list.

In [115...

```
he, she = get_occupation_stereotypes('she', 'he', occupations_file, word_to_vector_map)

print("Occupations closest to he:")
for occupation in he:
    print(f"{occupation[0], occupation[1]}")

print("\nOccupations closest to she:")
for occupation in she:
    if occupation[1] != 'businesswoman': # Excluding businesswoman
        print(f"{occupation[0], occupation[1]}")
```

Occupations closest to he:

```
(-0.35621238403428845, 'coach')
(-0.33694609670744985, 'caretaker')
(-0.316340865049864, 'captain')
(-0.30927324022117164, 'marshal')
(-0.30729135719953826, 'colonel')
(-0.3024871694158589, 'skipper')
(-0.3021438553042009, 'manager')
(-0.30165373167213716, 'midfielder')
(-0.2996711177716737, 'archbishop')
(-0.29443064216114034, 'commander')
(-0.2902867659336358, 'footballer')
(-0.2888985018624171, 'bishop')
(-0.2819996009929521, 'marksman')
(-0.2796388242138503, 'firebrand')
(-0.27878757562089734, 'provost')
(-0.2780793318566927, 'substitute')
(-0.27217972416807396, 'lieutenant')
(-0.2719793191353711, 'custodian')
(-0.2719191225303118, 'superintendent')
(-0.27134651130558013, 'goalkeeper')
```

Occupations closest to she:

```
(0.32072737162261483, 'singer')
(0.3219568449876504, 'publicist')
(0.3440519945920838, 'nanny')
(0.34466952852982563, 'therapist')
(0.3474650148960070, 'journalist')
```

```
(0.3474659148960079, 'confesses')
(0.3573048226883874, 'dancer')
(0.3645661876706785, 'hairstylist')
(0.36983542416768733, 'receptionist')
(0.3729106376588047, 'housekeeper')
(0.37309747916414077, 'homemaker')
(0.3812397049469023, 'housewife')
(0.3813398603461959, 'nurse')
(0.3892646060961875, 'narrator')
(0.4138336650968037, 'maid')
(0.42853139796990286, 'socialite')
(0.4434377286463442, 'waitress')
(0.44732913871667485, 'stylist')
(0.4686138536614218, 'ballerina')
(0.49840294304026295, 'actress')
```

## Task 4 - Debiasing word embeddings

### Gender Specific words

Words that are associated with a gender by definition. For example, brother, sister, businesswoman or businessman.

### Gender neutral words

The remaining words that are not specific to a gender are gender neutral. For example, flight attendant or shoes. The complement of gender specific words, can be taken as the gender neutral words.

### Step 1 - Identify gender subspace i.e identify the direction of the embedding that captures the bias

To robustly estimate bias, we use the gender specific words to learn a gender subspace in the embedding. To identify the gender subspace, we consider the vector difference of gender specific word pairs, such as  $\vec{she} - \vec{he}$ ,  $\vec{woman} - \vec{man}$  or  $\vec{her} - \vec{his}$ . This identifies a **gender direction or bias subspace**  $g \in \mathbb{R}^d$  which captures gender in the embedding.

**Note:** We will use  $g$  and *bias\_direction* interchangeably in this lab.

In [116...

```
gender = word_to_vector_map['she'] - word_to_vector_map['he']
print(gender)
```

```
[ 0.261302  0.438481 -0.13376   0.12281   0.00838   0.64455
  0.13151   0.01198   0.73557  -0.04754  -0.04261  -0.23387
  0.56951   0.24359   0.29471   0.152461 -0.44638   0.08563
  0.66735  -0.202578  0.28133   0.71557   0.04015   0.42204
  0.63574   0.1193   -0.429694  0.216301  0.08826  -0.5115
 -0.286     0.227249  0.25811   0.18075  -0.22733  -0.151844
 -0.13196  -0.14412   0.01709   -0.6281   0.124465  -0.16902
  0.6244667 -0.53734   0.379254  -0.3373   0.384876  -0.92383
 -0.019064  0.435641 ]
```

The gender subspace can also be captured more accurately by taking gender pair difference vectors and computing its principal components (PCs). The top PC, denoted by the unit vector  $g$ , captures the gender subspace.

In [117...

```
def get_gender_subspace(pairs, word_to_vector_map, num_components=10):
    """
    Compute the gender subspace by computing the principal components of
    ten gender pair vectors.
    Arguments:
        pairs (List[Tuple(String, String)]): A list of gender specific word pairs
        word_to_vector_map (Dict): A dictionary mapping words to embedding vectors
        num_components (Int): The number of principal components to compute. Defaults to 10
    Returns:
        gender_subspace (ndarray): The gender bias subspace(or direction) of shape (embedding dimension,)
    """

    matrix = []
    for word_1, word_2 in pairs:
        embedding_vector_word_1 = word_to_vector_map[word_1]
        embedding_vector_word_2 = word_to_vector_map[word_2]
        center = (embedding_vector_word_1 + embedding_vector_word_2) / 2
        matrix.append(embedding_vector_word_1 - center)
        matrix.append(embedding_vector_word_2 - center)

    matrix = np.array(matrix)
    pca = PCA(n_components=num_components)
    pca.fit(matrix)

    pcs = pca.components_           # Sorted by decreasing explained variance
    eigenvalues = pca.explained_variance_ # Eigenvalues
    gender_subspace = pcs[0]        # The first element has the highest eigenvalue
    return gender_subspace
```

In [118...

```
gender_specific_pairs = [
    ('she', 'he'),
    ('her', 'his'),
    ('woman', 'man'),
    ('mary', 'john'),
    ('herself', 'himself'),
    ('daughter', 'son'),
    ('mother', 'father'),
    ('gal', 'guy'),
    ('girl', 'boy'),
    ('female', 'male')
]
gender_direction = get_gender_subspace(gender_specific_pairs, word_to_vector_map)
print(gender_direction)
```

```
[ 0.06639123  0.15316702 -0.12170385  0.02910502 -0.012115  0.2956192
 0.10015248  0.03503806  0.27605339 -0.06259264  0.04843718 -0.20243709
 0.22435017  0.02205075  0.08795604  0.05350635 -0.23457441 -0.0051648
 0.29096997  0.02894429  0.10423079  0.24379617  0.05296573  0.17222571
 0.13557158  0.13746521 -0.05081975  0.11252051  0.01639264 -0.2113686
-0.1403471  0.13498117  0.08092433  0.02423979 -0.10780551 -0.05927322
-0.04857578 -0.03199024  0.08174041 -0.17759707 -0.02782478 -0.16880811
 0.27589146 -0.18007478  0.04123208 -0.09385728  0.11011447 -0.25650007
 0.06258361  0.00847159]
```

---

**Task 4a:** Run the cell below to compute the similarity between the gender embedding and the embedding vectors of male and female names. What can you observe?

---

In [119...

```
print('Names and their similarities with simple gender subspace')
names = ["mary", "john", "sweta", "david", "kazim", "angela"]
for name in names:
    print(name, cosine_similarity(word_to_vector_map[name], gender))

print()
print('Names and their similarities with PCA based gender subspace')
names = ["mary", "john", "sweta", "david", "kazim", "angela"]
for name in names:
    print(name, cosine_similarity(word_to_vector_map[name], gender_direction))
```

Names and their similarities with simple gender subspace

```
mary 0.3457399102816379
john -0.17879783833420468
sweta 0.17016456601128147
david -0.1332261560078667
kazim -0.32658964009764835
angela 0.2600799146632235
```

Names and their similarities with PCA based gender subspace

```
mary 0.2637091204419718
john -0.3816839789078354
sweta 0.1773704777691709
david -0.3165647635266187
kazim -0.3249838182709315
angela 0.18623308926276097
```

---

**Task 4b:** Quantify direct and indirect biases between words and the gender embedding by running the following cell. What is your observation?

---

In [120...

```
words = ["engineer", "science", "pilot", "technology", "lipstick", "arts", "singer", "computer", "receptionist", "fashion", ']
```

```

for word in words:
    print(word, cosine_similarity(word_to_vector_map[word], gender_direction))

```

```

engineer -0.2626286258749398
science -0.1202780958734538
pilot -0.1319833052724868
technology -0.1801116607819377
lipstick 0.4179404715417419
arts -0.04513818522820779
singer 0.16162975755073875
computer -0.16390549337211754
receptionist 0.3305284235998437
fashion 0.06913524872078784
doctor 0.02885191966409418
literature -0.08972688088254833

```

## Step 2 - Neutralize gender neutral words

Ensures that gender neutral words are zero in the gender subspace. This means that this steps takes a vector such as  $e_{fashion}$  and turns its components into zeros in the direction of  $g$  to produce  $e_{fashion}^{debiased}$

To remove bias in words such as "receptionist" or "shoe", given an input embedding of the word  $e$ , we compute debiased  $e$  denoted as  $e^{debiased}$  by using the formulas:

$$e^{bias\_component} = \frac{e \cdot bias\_direction}{||bias\_direction||_2^2} * bias\_direction \quad (3)$$

$$e^{debiased} = e - e^{bias\_component} \quad (4)$$

Where  $e^{bias\_component}$  is the projection of the word embedding  $e$  onto the gender subspace. Since the gender subspace is an orthogonal unit vector it is simply a direction. This also means that  $e^{debiased}$  is the projection onto the orthogonal subspace.

$||g||_2^2$  is the squared euclidean norm of  $g$  formulated as:

$$||g||_2^2 = \sum_i g_i^2$$

---

**Task 4c:** Implement `neutralize()` below by implementing the formulas above. Hint see [np.sum](#)

---

In [121...

```

def neutralize(word, gender_direction, word_to_vector_map):
    """
    Project the vector of word onto the gender subspace to remove the bias of "word"
    Arguments:
        word (String): A word to debias

```

```

gender_direction (ndarray): Numpy array of shape (embedding size (50), ) which is the bias axis
word_to_vector_map (Dict): A dictionary mapping words to embedding vectors

Returns:
    debiased_word (ndarray): the vector representation of the neutralized input word
"""

# Get the vector representation of word
embedding_of_word = word_to_vector_map[word]

# Compute the projection of word onto gender direction (eq. 3)
projection_of_word_onto_gender = (np.dot(embedding_of_word, gender_direction) / np.sum(gender_direction**2)) * gender_dir

# Neutralize word (eq. 4)
debiased_word = embedding_of_word - projection_of_word_onto_gender

return debiased_word

```

**Task 4d:** Test your implementation by running the code cell below. What is your observation?

In [122...

```

word = "babysit"
print(f"Before neutralization, cosine similarity between {word} and gender is: {cosine_similarity(word_to_vector_map[word], g

debiased_word = neutralize(word, gender_direction, word_to_vector_map)
print(f"After neutralization, cosine similarity between {word} and gender is: {cosine_similarity(debiased_word, gender_direct

```

Before neutralization, cosine similarity between babysit and gender is: 0.2663444879209918  
After neutralization, cosine similarity between babysit and gender is: -1.3389570015765782e-17

### Step 3 - Equalize

Equalizes sets of gender specific words outside the subspace. The goal is to ensure that gender neutral words are equidistance to all the words in the set. We want to ensure that gender specific words are not biased with respect to neutral words.

For example, consider the set {woman, man}, if the neutral word "babysit" is closer to "woman" than "man" then the neutralization of "babysit" can reduce the gender-stereotype associated with babysitting but does not make "babysit" equidistant to "woman" and "man".

Given two gender specific word pairs  $w_1$  and  $w_2$  to debias, and their embeddings  $e_{w_1}$  and  $e_{w_2}$ , equalization can be achieved with the following equations:

$$\mu = \frac{e_{w_1} + e_{w_2}}{2} \quad (5)$$

$$\mu_B = \frac{\mu \cdot \text{bias\_direction}}{\|\text{bias\_direction}\|_2^2} * \text{bias\_direction} \quad (6)$$

$$v = \mu - \mu_B \quad (7)$$

$$e_{w_1B} = \frac{e_{w_1} \cdot \text{bias\_direction}}{\|\text{bias\_direction}\|_2^2} * \text{bias\_direction} \quad (8)$$

$$e_{w_2B} = \frac{e_{w_2} \cdot \text{bias\_direction}}{\|\text{bias\_direction}\|_2^2} * \text{bias\_direction} \quad (9)$$

$$e_{w_1B}^{new} = \sqrt{|1 - \|v\|_2^2|} * \frac{e_{w_1B} - \mu_B}{\|(e_{w_1} - v) - \mu_B\|_2} \quad (10)$$

$$e_{w_2B}^{new} = \sqrt{|1 - \|v\|_2^2|} * \frac{e_{w_2B} - \mu_B}{\|(e_{w_2} - v) - \mu_B\|_2} \quad (11)$$

$$e_1 = v + e_{w_1B}^{new} \quad (12)$$

$$e_2 = v + e_{w_2B}^{new} \quad (13)$$

---

**Task 5a:** Implement `equalization()` below by implementing the formulas above.

---

In [123...

```
def equalization(equality_set, bias_direction, word_to_vector_map):
    """
    Equalize the pair of gender specific words in the equality set ensuring that
    any neutral word is equidistant to all words in the equality set.
    Arguments:
        equality_set (Tuple(String, String)): a tuple of strings of gender specific
        words to debias e.g ("grandmother", "grandfather")
        bias_direction (ndarray): numpy array of shape (embedding dimension,). The
        embedding vector representing the bias direction
        word_to_vector_map (Dict): A dictionary mapping words to embedding vectors
    Returns:
        embedding_word_a (ndarray): numpy array of shape (embedding dimension,). The
        embedding vector representing the first word
        embedding_word_b (ndarray): numpy array of shape (embedding dimension,). The
        embedding vector representing the second word
    """

    # Start code here #
    # Get the vector representation of word pair by unpacking equality_set (~ 3 line)
    word_a, word_b = None
    embedding_word_a = None
```

```
embedding_word_b = None
```

```
# Compute the mean (eq. 5) of embedding_word_a and embedding_word_a (~ 1 line)
```

```
mean = None
```

```
# Compute the projection of mean representation onto the bias direction (eq. 6) (~ 1 line)
```

```
mean_B = None
```

```
# Compute the projection onto the orthogonal subspace (eq. 7) (~ 1 line)
```

```
mean_orthogonal = None
```

```
# Compute the projection of th embedding of word a onto the bias direction (eq. 8) (~ 1 line)
```

```
embedding_word_a_on_bias_direction = None
```

```
# Compute the projection of th embedding of word b onto the bias direction (eq. 9) (~ 1 line)
```

DA6233ECruz--2023 / IS6733Lab4COLABECruzkml614.ipynb

↑ Top

Preview

Code

Blame

1 lines (1 loc) · 51.9 KB

Raw



```
new_embedding_word_b_on_bias_direction = None
```

```
# Equalize embedding of word a using eq. 12 (~ 1 line)
```

```
embedding_word_a = None
```

```
# Equalize embedding of word b using eq. 13 (~ 1 line)
```

```
embedding_word_b = None
```

```
# End code here #
```

```
return embedding_word_a, embedding_word_b
```

**Task 5b:** Test your implementation by running the cell below.

In [124...

```
def equalization(equality_set, bias_direction, word_to_vector_map):
```

```
    """
```

```
    Equalizes the pair of gender-specific words in the equality_set ensuring that  
    any neutral word is equidistant to all words in the equality_set.
```

```
    Arguments:
```

```
        equality_set (Tuple(String, String)): a tuple of strings of gender-specific  
        words to debias e.g ("grandmother", "grandfather")
```

```
        bias_direction (ndarray): numpy array of shape (embedding_dimension,). The  
        embedding vector representing the bias direction
```

```
        word_to_vector_map (Dict): A dictionary mapping words to embedding vectors
```



```

Returns:
    embedding_word_a (ndarray): numpy array of shape (embedding_dimension,). The
        embedding vector representing the first word
    embedding_word_b (ndarray): numpy array of shape (embedding_dimension,). The
        embedding vector representing the second word
"""

# Get the vector representation of word pair by unpacking equality_set
word_a, word_b = equality_set
embedding_word_a = word_to_vector_map[word_a]
embedding_word_b = word_to_vector_map[word_b]

# Add this line to return the embedding vectors
return embedding_word_a, embedding_word_b

print("Cosine similarity before equalization:")
print(f"(embedding vector of father, gender_direction): {cosine_similarity(word_to_vector_map['father'], gender_direction)}")
print(f"(embedding vector of mother, gender_direction): {cosine_similarity(word_to_vector_map['mother'], gender_direction)}")
print()

embedding_word_a, embedding_word_b = equalization(("father", "mother"), gender_direction, word_to_vector_map)
print("Cosine similarity after equalization:")
print(f"(embedding vector of father, gender_direction): {cosine_similarity(embedding_word_a, gender_direction)}")
print(f"(embedding vector of mother, gender_direction): {cosine_similarity(embedding_word_b, gender_direction)}")

```

Cosine similarity before equalization:  
 (embedding vector of father, gender\_direction): -0.08502503175882657  
 (embedding vector of mother, gender\_direction): 0.3332593015356538

Cosine similarity after equalization:  
 (embedding vector of father, gender\_direction): -0.08502503175882657  
 (embedding vector of mother, gender\_direction): 0.3332593015356538

---

**Task 5c:** Looking at the output of your implementation test above, what can you observe?.

---

#### References:

- The debiasing algorithm is from Bolukbasi et al., 2016 [Man is to Computer Programmer as Woman is to Homemaker? Debiasing word Embeddings](#)
- The code is partly adapted from Andrew Ng's debiasing word embeddings course on [Coursera](#)
- The GloVe word embeddings is publicly available at (<https://nlp.stanford.edu/projects/glove/>) and is due to the works of Jeffrey Pennington, Richard Socher, and Christopher D. Manning.

